


Mean cost ranking

Mean API cost per response across 24 common exercises | top to bottom: cheaper -> more expensive


HOW TO READ  Dot = mean cost; horizontal line with end caps = 95% confidence interval



Axis starts at zero. Gemma 4 31B IT uses the campaign-configured zero rate; this is not a universal model price.

Overall token composition

Tokens accumulated across 24 common exercises | ordered from lower to higher volume



Claude Fable 5 does not report hidden thinking separately: the purple component is intentionally left unsplit.

How mean cost is formed


Consumed volume and effective unit cost jointly explain cost per response | each bubble is one model



The effective rate includes the observed input/output mix, so it need not equal either published list price.

Rates per million tokens

Unit prices used in the benchmark | direct comparison, independent of consumed volume



Reasoning is billed as output where applicable. Gemma uses the campaign-configured zero rate.