

Technical specifications of the evaluated models

A priori comparison of the 14 benchmark configurations. No benchmark outcome is included. Token limits and model status refer to first-party documentation available on 2026-08-10.

Model / API identifier	Public release	Parameters	Open weights	Architecture / model class	Reasoning capability / benchmark mode	Context window	Maximum output	Knowledge / training cutoff
GPT-5.6 Sol OPENAI gpt-5.6-sol	9 Jul 2026 GA; preview 26 Jun 2026	Not disclosed	No	Closed frontier flagship; architecture undisclosed	Configurable effort: none to max; benchmark: xhigh	1,050,000	128,000	16 Feb 2026 knowledge
OpenAI o3 OPENAI o3-2025-04-16	16 Apr 2025 dated snapshot	Not disclosed	No	Closed o-series reasoning model; architecture undisclosed	Native inference-time reasoning; benchmark: high	200,000	100,000	1 Jun 2024 knowledge
OpenAI GPT-4.1 OPENAI gpt-4.1-2025-04-14	14 Apr 2025 dated snapshot	Not disclosed	No	Closed general-purpose non-reasoning model	No native reasoning step; benchmark: none	1,047,576	32,768	1 Jun 2024 knowledge
Claude Fable 5 ANTHROPIC claude-fable-5	9 Jun 2026 GA; restored 1 Jul 2026	Not disclosed	No	Closed Mythos-class safeguarded model	Always-on adaptive thinking; benchmark: effort high	1,000,000	128,000	Jan 2026 reliable/training
Claude Sonnet 4.5 ANTHROPIC claude-sonnet-4-5-20250929	29 Sep 2025 dated snapshot	Not disclosed	No	Closed Sonnet-tier hybrid reasoning model	Optional manual extended thinking; benchmark: provider default	200,000	64,000	Not disclosed
Claude Haiku 4.5 ANTHROPIC claude-haiku-4-5-20251001	15 Oct 2025 API snapshot 1 Oct 2025	Not disclosed	No	Closed small and low-latency tier	Manual extended thinking, off by default; benchmark: not requested	200,000	64,000	Feb/Jul 2025 reliable/training
Gemini 3.5 Flash GOOGLE gemini-3.5-flash	19 May 2026 GA	Not disclosed	No	Closed multimodal Flash-tier model	Configurable native thinking; benchmark: high	1,048,576	65,536	Jan 2025 knowledge
Gemini 3.1 Pro Preview GOOGLE gemini-3.1-pro-preview	19 Feb 2026 public preview	Not disclosed	No	Closed multimodal Pro-tier model	Configurable native thinking; benchmark: high	1,048,576	65,536	Jan 2025 knowledge
Gemma 4 31B IT GOOGLE gemma-4-31b-it	31 Mar 2026 API 2 Apr 2026	30.7B total; dense; 60 layers; about 550M vision encoder	Yes; Apache 2.0	Dense decoder-only transformer; hybrid local/global attention	Configurable thinking tokens; benchmark: disabled	256,000	Not specified	Jan 2025 training
DeepSeek V4 Flash DEEPSEEK deepseek-v4-flash - thinking	24 Apr 2026 V4 preview	284B total; 13B active per token	Yes	Mixture-of-Experts; DeepSeek Sparse Attention	Dual thinking/non-thinking model; benchmark: thinking high	1,000,000	384,000	Not disclosed
DeepSeek V4 Flash (no thinking) DEEPSEEK deepseek-v4-flash - no thinking	24 Apr 2026 same weights as above	284B total; 13B active per token	Yes	Mixture-of-Experts; DeepSeek Sparse Attention	Dual thinking/non-thinking model; benchmark: disabled	1,000,000	384,000	Not disclosed
DeepSeek V4 Pro DEEPSEEK deepseek-v4-pro	24 Apr 2026 V4 preview	1.6T total; 49B active per token	Yes	Mixture-of-Experts; DeepSeek Sparse Attention	Dual thinking/non-thinking model; benchmark: thinking high	1,000,000	384,000	Not disclosed
Grok 4.5 XAI grok-4.5	16 Jul 2026 public launch	Not disclosed	No	Closed frontier model; architecture undisclosed	Configurable low/medium/high; benchmark: high	500,000	Not specified	1 Feb 2026 knowledge
Grok 4.3 XAI grok-4.3	By 6 May 2026* first dated official use	Not disclosed	No	Closed general-purpose model; architecture undisclosed	Configurable none/low/medium/high; benchmark: high	1,000,000	Not specified	Not disclosed

* xAI provides no dedicated dated launch announcement for Grok 4.3; 6 May 2026 is the earliest dated official production reference identified. 'Not disclosed' is used when the developer publishes no official value; no third-party parameter estimates are included. The two V4 Flash rows are separate benchmark configurations of the same model weights.

Primary sources: OpenAI model documentation; Anthropic model overview and release notes; Gemini API model and deprecation documentation; Gemma 4 model card; DeepSeek V4 release and pricing documentation; xAI model documentation.